

## RESEARCH ARTICLE

# Large scale paired antibody language models

Henry Kenlay<sup>1</sup>✉, Frédéric A. Dreyer<sup>1</sup>\*, Aleksandr Kovaltsuk<sup>1</sup>, Dom Miketa<sup>1</sup>, Douglas Pires<sup>1</sup>, Charlotte M. Deane<sup>1,2</sup>

**1** Exscientia, Oxford Science Park, Oxford, United Kingdom, **2** Department of Statistics, University of Oxford, Oxford, United Kingdom

✉ These authors contributed equally to this work.

\* [dreyer.frederic@gene.com](mailto:dreyer.frederic@gene.com)

**OPEN ACCESS**

**Citation:** Kenlay H, Dreyer FA, Kovaltsuk A, Miketa D, Pires D, Deane CM (2024) Large scale paired antibody language models. PLoS Comput Biol 20(12): e1012646. <https://doi.org/10.1371/journal.pcbi.1012646>

**Editor:** Miguel Francisco de Almeida Pereira de Rocha, University of Minho, PORTUGAL

**Received:** June 30, 2024

**Accepted:** November 18, 2024

**Published:** December 6, 2024

**Peer Review History:** PLOS recognizes the benefits of transparency in the peer review process; therefore, we enable the publication of all of the content of peer review and author responses alongside final, published articles. The editorial history of this article is available here: <https://doi.org/10.1371/journal.pcbi.1012646>

**Copyright:** © 2024 Kenlay et al. This is an open access article distributed under the terms of the [Creative Commons Attribution License](#), which permits unrestricted use, distribution, and reproduction in any medium, provided the original author and source are credited.

**Data Availability Statement:** The original studies and associated datasets for all sequence data used to train the model can be found at <https://opig.stats.ox.ac.uk/webapps/oas/>, with the list of the unpaired data accessible at <https://opig.stats.ox.ac.uk>.

## Abstract

Antibodies are proteins produced by the immune system that can identify and neutralise a wide variety of antigens with high specificity and affinity, and constitute the most successful class of biotherapeutics. With the advent of next-generation sequencing, billions of antibody sequences have been collected in recent years, though their application in the design of better therapeutics has been constrained by the sheer volume and complexity of the data. To address this challenge, we present IgBert and IgT5, the best performing antibody-specific language models developed to date which can consistently handle both paired and unpaired variable region sequences as input. These models are trained comprehensively using the more than two billion unpaired sequences and two million paired sequences of light and heavy chains present in the Observed Antibody Space dataset. We show that our models outperform existing antibody and protein language models on a diverse range of design and regression tasks relevant to antibody engineering. This advancement marks a significant leap forward in leveraging machine learning, large scale data sets and high-performance computing for enhancing antibody design for therapeutic development.

## Author summary

This study introduces IgBert and IgT5, two advanced antibody-specific language models capable of processing the extensive data generated by next-generation sequencing of antibodies. These models are trained on over two billion unpaired and two million paired antibody sequences from the Observed Antibody Space dataset. IgBert and IgT5 demonstrate state-of-the-art performance compared to existing models in tasks such as sequence recovery, expression, and binding affinity predictions. By employing comprehensive machine learning techniques and leveraging large-scale datasets, these models significantly enhance the potential for antibody design and therapeutic development. The trained models are publicly available, providing valuable tools for ongoing research in antibody engineering.

[uk/webapps/oas/oas\\_unpaired/](https://uk/webapps/oas/oas_unpaired/), and the list of paired data accessible at [https://opig.stats.ox.ac.uk/webapps/oas/oas\\_paired/](https://opig.stats.ox.ac.uk/webapps/oas/oas_paired/). To reach the original studies and associated datasets, users need to click the "search" button with no filters. The models weights are available on Zenodo at <https://doi.org/10.5281/zenodo.10876909>. The code used is available through the HuggingFace library at <https://huggingface.co/exscientia>.

**Funding:** The author(s) received no specific funding for this work.

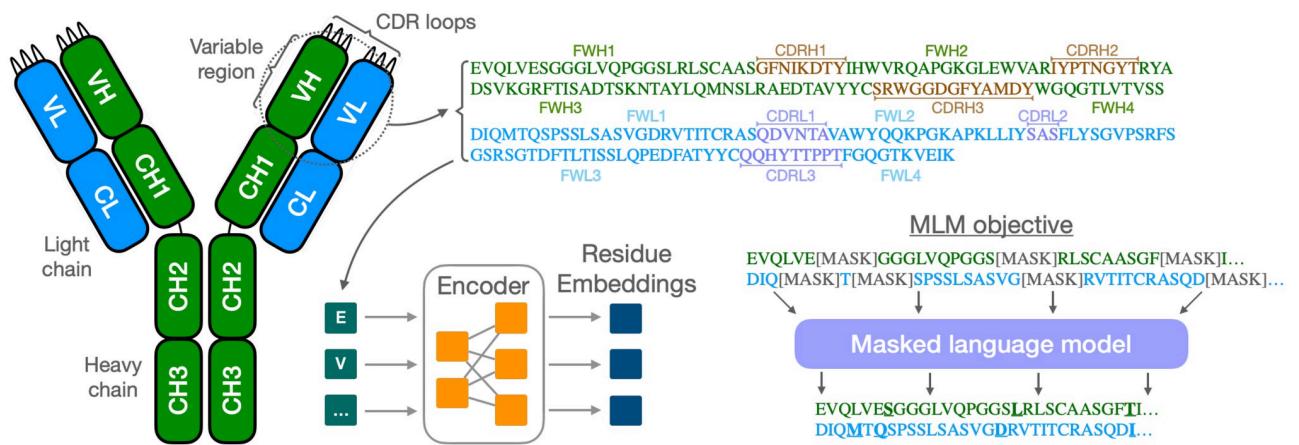
**Competing interests:** The authors have declared that no competing interests exist.

## Introduction

Antibodies are proteins that play a central role in the adaptive immune system recognising and neutralising pathogens. They are symmetric Y-shaped molecules consisting of two light and two heavy chains (as depicted in Fig 1), and can identify and bind with high specificity to antigens—substances that the body identifies as foreign, such as bacteria, viruses, or toxins. Antibodies possess a variable region at their tip that allows them to recognise antigens, which is one of the ways they activate further immune functions. An antibody's specificity is primarily influenced by the complementarity-determining regions (CDRs) of the variable domain. The third CDR loop of the heavy chain is the most structurally diverse and typically the most important for antigen recognition [1–3]. The ability to specifically target pathogens makes antibodies a central element in both natural immune responses and medical applications.

The emergence of next-generation sequencing (NGS) technology has transformed our understanding of the diversity and function of antibodies [4]. NGS enables the rapid sequencing of antibody repertoires, revealing a detailed insight into the natural diversity of the variable region of antibodies. This extensive sequencing data provides a large-scale representation of the various antibodies that natural systems generate, and contains crucial information about the numerous ways in which antibodies can adapt to recognise a wide array of antigens. This data is key to understanding how the immune system responds to infections and how it can be manipulated for therapeutic purposes.

The vast amount of data generated by antibody sequencing presents a unique challenge, which traditional bioinformatics methods such as sequence alignment and motif search can only partially address. Protein language models, inspired by advances in natural language processing, offer an orthogonal approach, which have found a wide range of uses in protein design by learning rich representations from large corpora of unlabelled sequences. They can be trained specifically to learn the complex grammar of antibody sequences, often through a masked language model objective, where a random fraction of amino acid residues are hidden and the model is then tasked to reconstruct the original sequence. This unsupervised pre-training allows the model to learn a high dimensional contextualised embedding space on which sequences can be mapped, enabling a wide range of applications such as restoring missing residues [5], *de novo* sequence generation [6, 7], structure [8, 9] and property prediction such as binding sites [10], humanness [11] or thermostability [12]. Language models thus represent a



**Fig 1. Overview of an antibody structure and its domains.** The sequence of the variable region is used as input to the transformer encoder to obtain a residue-level embedding representation. Training is achieved through masked language modelling, where a random fraction of the input is replaced by mask tokens.

<https://doi.org/10.1371/journal.pcbi.1012646.g001>

crucial tool in translating the vast antibody sequencing data into insights that can be used to advance therapeutic developments.

In this work, we train antibody language models on the Observed Antibody Space (OAS) [13, 14]. We consider both a BERT [15] and a T5 [16] model, which are initialised with weights of a general protein model and pre-trained on a large corpus of unpaired antibody sequences, consisting of the totality of known unpaired sequences for the BERT model, and 700M cluster representatives for the T5 model. These models are then fine-tuned on a smaller dataset of all known paired heavy and light variable region sequences. To our knowledge, these are the largest and one of the only paired antibody language models, as well as the first T5 antibody language model trained for sequence encoding to date. We then show how our models perform on downstream applications, notably sequence recovery, as well as expression and binding affinity predictions, and demonstrate that they outperform existing state-of-the-art protein and antibody language models on these key tasks.

The models trained as part of this study are made publicly available and can be readily used in a range of tasks of relevance to antibody engineering [17], as detailed in [S1 Appendix](#).

## Related work

The vast amount of sequence data available in public repositories positions unsupervised representation learning approaches as powerful ways to leverage machine learning in biology. Protein language models, notably transformer-based architectures such as Evolutionary Scale Modeling (ESM) [18, 19], ProtTrans [20], and ProGen [21] have been pivotal in protein structure prediction and function annotation, providing a solid foundation for generalised protein modeling. Trained on extensive evolutionary data containing billions of sequences [22, 23], these models have also demonstrated remarkable capabilities in understanding protein sequences, their evolutionary trajectories and constraints [24].

Turning to antibody-specific literature, BERT-like models [15, 25] such as AbLang [5], AntiBERTy [26], Sapiens [11], and AntiBERTa [10] have shown that training language models on sequences from the immunoglobulin protein superfamily can be of particular use, notably to restore missing residues, to study affinity maturation trajectories and to identify paratope residues (*i.e.*, those responsible for recognising antigen binding). Models trained with paired sequences of heavy and light chains can learn cross-chain features, which has been shown to improve performance on predictive downstream tasks [27, 28]. Effective antibody-specific models can also be obtained by fine-tuning general protein models on antibody sequences [27, 29], which can often be less computationally intensive than training a model from random initial weights.

While most antibody language models have focused on encoder-only architectures, some generative models have explored encoder-decoder [30] and decoder-only models [6, 31].

## Methods

### Models

In this study, we focus on two language model architectures, BERT [15] and T5 [16]. BERT is an encoder-only model focused on generating contextual embeddings for each token in an input sequence whereas T5 is an encoder-decoder model designed for sequence-to-sequence tasks, where the encoder processes the input sequence and the decoder generates the output sequence. Both are based on transformer architectures [32], where inputs are processed through a series of blocks that incorporate self-attention with feed-forward connections. The self-attention component allows the network to construct representations that incorporate

context from anywhere in the sequence, and which can directly encode pairwise interactions between amino acid residues.

Bidirectional Encoder Representations from Transformers (BERT) is a bidirectional model which learns a contextual representation by predicting randomly masked tokens in the input. By incorporating information from both left and right context across each layer, it can capture nuanced understanding of language structure and context and has seen diverse applications in natural language processing. It is typically pre-trained on a large corpus of text, which in our case are protein and antibody sequences, and then fine-tuned for specific tasks, including but not limited to sentiment analysis, language translation, or in the case of proteins, structural and functional property prediction.

Text-To-Text Transfer Transformer (T5) frames all tasks as text-to-text problems, and employs span-based masking where contiguous random spans are replaced with sentinel tokens. The decoder uses a causal attention mechanism, which allows it to handle a diverse range of tasks including translation, summarisation and question answering with a unified approach.

We train both models with a masked language modelling (MLM) objective. Sequences are parsed through a tokeniser which maps each amino acid type in the sequence to a distinct token, and includes a separator token at the end of each chain. The BERT tokeniser has a classification and mask token, while the T5 tokeniser has a single sentinel token which is used as an equivalent to a mask token. During training, each input sequence is corrupted by replacing a fraction of the amino acid residues with a special mask token for the BERT model, while for the T5 model we follow [20], with a corresponding single-token span. The models are then trained to predict these tokens from the corrupted sequence, such that for a sequence with  $n$  residues  $\mathbf{s} = (s_1, \dots, s_n)$  the loss is defined as

$$\mathcal{L}_{\text{MLM}} = \mathbb{E} \left[ \sum_{i \in \mathcal{M}} -\log p(s_i | \mathbf{s}_{/\mathcal{M}}) \right], \quad (1)$$

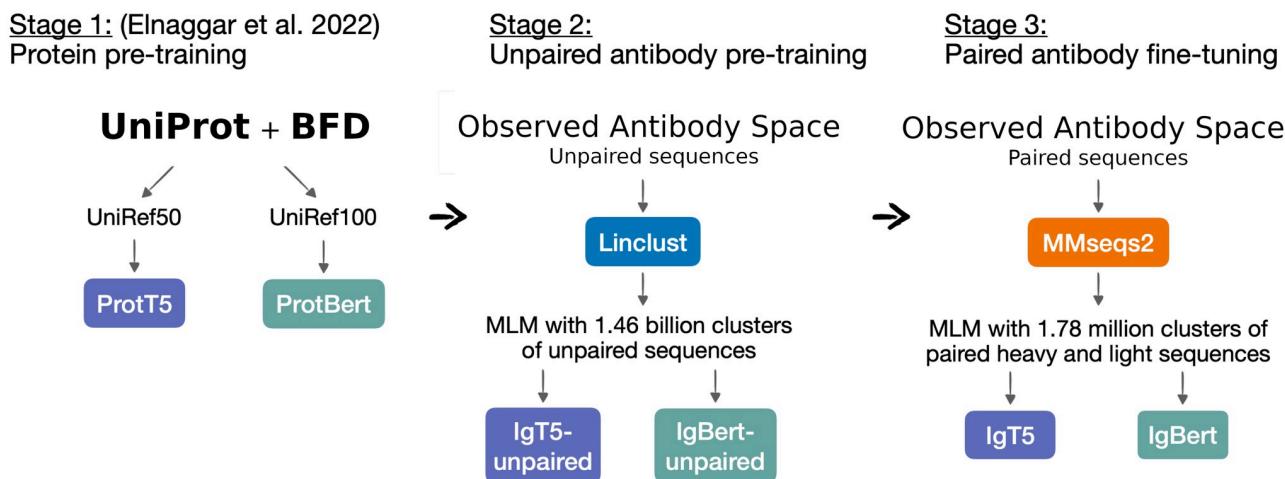
where  $\mathcal{M}$  are all the masked positions in a sequence and  $\mathbf{s}_{/\mathcal{M}}$  denotes the masked sequence where all masked positions have been replaced by a mask token.

A key difference between the BERT and T5 models on this MLM task is that the BERT model has a simple MLM head to predict the token based on the encoded embedding, while the T5 model is trained to predict the sequence as a text-to-text task and relies on both the encoder and decoder, even if only the encoder is later used in most downstream tasks. The loss is then computed over the full sequence for the T5 model, while only the masked residues contribute to the BERT loss.

For both the T5 and BERT model, we train our models starting from existing weights of a general protein language model that was pre-trained on protein sequences from the UniRef [33] and BFD [34] databases. We use as starting point the best performing BERT-style model, ProtBert, and the state-of-the-art ProtT5 model, both from the ProtTrans family of protein transformer models [20]. We found that this approach generally achieves higher performance for a fixed compute budget than training from random weights, as was similarly reported in [27].

## Data preparation

We train our models on sequences from OAS, a database of predominantly human immune repertoires which contains over two billion unpaired variable heavy and light chain sequences, as well as over two million paired variable region sequences. For a brief summary of the OAS dataset, see S2 Appendix. An overview of our training strategy is shown in Fig 2, which



**Fig 2. Data processing and training strategy.** We further pre-train the ProtT5 and ProtBert models from [20] on unpaired antibody sequences from OAS after clustering them with Linclust. These unpaired models are then fine-tuned on paired sequences clustered with MMseqs2, combining the two variable region chains into a single input with a separator token.

<https://doi.org/10.1371/journal.pcbi.1012646.g002>

includes the initial pre-training on general protein sequences from UniRef [33] and BFD [34] and initial weights from the ProtTrans models [20].

We perform a second stage pre-training on the unpaired OAS sequences. We take all 2,097,593,973 unique sequences from the database after removing duplicates. We cluster the sequences to ensure no sequences in the validation and test sets are near-duplicates to those in the training set. Due to the large size of this dataset, we cluster these sequences with Linclust [35], an approximate clustering algorithm that scales linearly with the number of sequences. We use 95% sequence identity as a threshold to create a dataset of 1,456,516,479 clusters. From these clusters, 40M are randomly selected to construct test and validation clusters. The remaining 1,416,516,479 clusters compose the training dataset, which contains 2,040,122,161 unique sequences. The validation and test datasets consist of 20M representative sequences each, selected from the 40M clusters selected above.

The last stage is fine-tuning on paired sequences. Here we consider the 2,038,528 unique paired antibody variable region sequences from OAS, and cluster them with MMseqs2 [36] using a 95% sequence identity threshold on the concatenated heavy and light chains. In contrast to Linclust, MMseqs2 is an exact clustering algorithm which ensures non-overlapping clusters, but scales quadratically with the number of sequences. This provides us with a dataset of 1,777,722 clusters, of which 40,000 are randomly selected as test and validation clusters. The training dataset consists of 1,992,786 unique sequences, distributed among 1,737,722 clusters. As validation and test data, we use the 20,000 representative sequences drawn from the test and validation clusters.

### Pre-training on unpaired sequences

All models were trained using 32 A100 GPUs using a distributed DeepSpeed ZeRO Stage 2 strategy, which shards optimiser states and gradients for memory efficiency [37]. We fix the number of training steps ahead of time and linearly warm-up the learning rate to the maximum learning rate before linearly annealing to zero. The optimiser for both models is AdamW with a weight decay parameter of  $10^{-5}$  for regularisation [38]. Each sequence was padded to a length of 200. We summarise the model and training parameters for all models in Table 1.

**Table 1. Hyperparameters for all models trained in this study.** The IgBert-unpaired and IgT5-unpaired are initialised with the ProtBert and ProtT5 weights respectively. The IgBert model is trained starting from the IgBert-unpaired weights, and the IgT5 model is derived from IgT5-unpaired.

Parameters	IgBert-unpaired	IgT5-unpaired	IgBert	IgT5
Initial weights	ProtBert	ProtT5	IgBert-unpaired	IgT5-unpaired
Number of layers	30	24	30	24
Size hidden layers	1024	1024	1024	1024
Intermediate size hidden layers	4096	16,384	4096	16,384
Number of attention heads	16	32	16	32
Dropout	0.0	0.1	0.0	0.1
Target length	200	200	300	300
Batch size per GPU	300	24	192	15
Global batch size	9600	768	6144	480
Optimiser	AdamW	AdamW	AdamW	AdamW
Max learning rate	$10^{-5}$	$5 \cdot 10^{-5}$	$10^{-5}$	$5 \cdot 10^{-5}$
Weight decay	$10^{-5}$	$10^{-5}$	$10^{-5}$	$10^{-5}$
Warm-up steps	10,000	30,000	2700	4000
Total training steps	212,512	911,500	45,000	60,000
Masking probability	0.15	0.15	0.15	0.15
Number of parameters	420M	3B	420M	3B

<https://doi.org/10.1371/journal.pcbi.1012646.t001>

**IgBert-unpaired.** The IgBert-unpaired model is initialised with ProtBert weights. We fine-tune on all unpaired OAS sequences that belong to training clusters. The batch size is set to 300 per GPU for a global batch size of 9600.

To complete a full epoch we trained for 212,512 steps. We used a linear annealing scheduler with a max learning rate of  $10^{-5}$  and 10,000 warm-up steps. Training time was approximately 66 hours.

We used a MLM objective where 15% of tokens were selected for masking. Of these 80% are replaced by a masked token, 10% are changed to a random token in the vocabulary and 10% are left unchanged.

**IgT5-unpaired.** The T5 architecture is much larger than the BERT architecture in terms of parameters leading to a larger memory and compute footprint. As such, we train on just representative sequences of the training clusters for half an epoch due to compute constraints. Similar to before, we initialise training using the ProtT5 weights.

The batch size is 24 per GPU for a global batch size of 768. We trained for a total of 911,500 steps, which corresponds to 700M representative sequences or half of the unpaired cluster representatives. The max learning rate was set to a larger  $5 \times 10^{-5}$ , in line with the proportionally larger learning rate used to train ProtT5 compared to ProtBert [20]. The warm-up was 30,000 steps. Training time was approximately 9 days.

Following ProtT5, we used the T5 objective with a maximum span length of one. This leads to a similar loss to the MLM objective, but with an additional sequence reconstruction term. We masked 15% of tokens where 90% are replaced by a mask token and 10% are replaced by a random amino acid token.

### Fine-tuning on paired sequences

We further tune our unpaired models on paired sequences available in the OAS. To prevent catastrophic forgetting [39, 40], we use batches with both unpaired and paired data, incorporating two unpaired sequences into the batch for each paired sequence included. Paired

sequence inputs are constructed by concatenating the heavy and light chain sequences, inserting a separator token in between them.

The input sequences were padded to 300 to account for the increased length of paired sequences, reducing the batch size we could fit into memory for both models. Apart from the learning rate schedule and batch construction the rest of the setup remained unchanged. We fine-tuned both models for 16 hours.

**IgBert.** IgBert was trained by fine-tuning IgBert-unpaired using a batch size of 128 unpaired and 64 paired sequences per GPU. The paired training data consists of all sequences in the training clusters for the paired sequences, and randomly chosen unpaired sequences from the unpaired training clusters to complete the batches. We trained for approximately 46 epochs (45,000 steps), defining an epoch as one pass over the paired OAS dataset. The warm-up consisted of 2,700 steps, approximately three epochs. We use the same MLM objective as in the unpaired pre-training, where 15% of the input is masked, 80% of those are replaced by a mask token, 10% by a random token and 10% are kept identical.

**IgT5.** Similar to IgBert, we used IgT5-unpaired as our starting weights to train IgT5. The batch size was 10 unpaired and 5 paired sequences per GPU. The batches are constructed from all sequences from the paired training clusters, and representative sequences randomly selected from the unpaired training clusters which have not been seen in the unpaired pre-training. We trained for approximately 5 epochs (60,000 steps) of the paired OAS data, with 4,000 warm-up steps. The objective is identical to the unpaired pre-training case, with 15% masked tokens, from which 90% are replaced by a sentinel token and 10% are randomly replaced by a different amino acid.

## Results

We evaluate our models on a range of downstream tasks, comparing with existing antibody and protein language models. As baselines, we consider the original ProtTrans protein language models [20], as well as the open-source antibody language models AbLang [5] and AntiBERTy [26]. We examine two tasks of interest: the sequence recovery for each component of the antibody variable region, and performance in predicting experimental binding affinity or expression data from language model embeddings.

### Sequence recovery

In Table 2, we show the sequence recovery achieved after masking 15% of residues in paired sequences from the test dataset at random. Similar results for unpaired sequences are provided in S3 Appendix. The recovery fraction is shown separately for each framework and CDR loop, as well as over the total heavy and light sequences.

This task benefits from training on antibody-specific data, as can be seen by the comparatively poor performance of the protein language models, notably in the hypervariable H3 and L3 loops. We also observe that the IgT5 model achieves the highest recovery across all but one of the CDR loops, CDRL1, where IgBert achieves marginally higher accuracy. For sequence recovery in the framework regions, both IgBert and IgT5 have comparable high performance.

### Binding affinity and expression

We now consider the language model embeddings as input feature to a predictive downstream task. To this end, we consider the largest datasets from the FLab benchmark [41], which include the binding energy datasets from [42] with 422 data points, from [43] with 2048 data points and from [44] with 4275 data points, as well as the corresponding expression data from this last study.

**Table 2.** Fraction of correctly predicted residues by region (frameworks (FW) and CDRs for the heavy (H) and light (L) chains), after masking 15% of the sequence for a test set of 20k paired heavy and light sequences. The best, second and third best performing models for each region are shown in bold, underlined and italic respectively.

Model	FWH1	FWH2	FWH3	FWH4	CDRH1	CDRH2	CDRH3	Total VH
AbLang [5]	0.979	0.967	0.956	0.981	0.910	0.884	0.593	0.910
AntiBERTy [26]	0.978	0.965	0.954	0.977	0.907	0.882	0.521	0.900
ProtBert [20]	0.802	0.761	0.738	0.846	0.656	0.456	0.277	0.682
IgBert-unpaired	0.979	0.965	0.955	0.980	0.904	0.884	0.592	0.910
IgBert	<b>0.981</b>	<b>0.969</b>	<b>0.958</b>	<b>0.981</b>	<u>0.913</u>	<u>0.886</u>	<u>0.601</u>	<u>0.913</u>
ProtT5 [20]	0.904	0.854	0.888	0.914	0.753	0.629	0.339	0.793
IgT5-unpaired	0.979	0.967	0.956	<u>0.982</u>	0.909	0.884	<u>0.603</u>	<u>0.912</u>
IgT5	<b>0.982</b>	<u>0.969</u>	<u>0.957</u>	<b>0.983</b>	<b>0.915</b>	<b>0.894</b>	<b>0.620</b>	<b>0.916</b>
Model	FWL1	FWL2	FWL3	FWL4	CDRL1	CDRL2	CDRL3	Total VL
AbLang [5]	0.966	0.968	0.971	0.962	0.891	0.901	0.838	0.949
AntiBERTy [26]	0.979	0.969	0.975	0.966	0.907	0.895	0.844	0.955
ProtBert [20]	0.660	0.786	0.783	0.634	0.469	0.438	0.290	0.665
IgBert-unpaired	0.980	0.970	0.974	0.966	0.908	0.898	0.846	0.956
IgBert	<b>0.988</b>	<b>0.974</b>	<u>0.981</u>	<u>0.974</u>	<b>0.923</b>	<u>0.915</u>	<u>0.863</u>	<u>0.965</u>
ProtT5 [20]	0.846	0.901	0.880	0.850	0.696	0.604	0.517	0.820
IgT5-unpaired	0.981	0.967	0.975	0.917	0.908	0.909	0.842	0.951
IgT5	<u>0.988</u>	<u>0.973</u>	<b>0.981</b>	<b>0.978</b>	<u>0.922</u>	<b>0.916</b>	<b>0.869</b>	<b>0.966</b>

<https://doi.org/10.1371/journal.pcbi.1012646.t002>

For each dataset, we use a linear least squares fit with  $L_2$  regularisation on the embedding representation features assessed under 10-fold cross-validation. We use a 5-fold inner cross-validation to select the regularisation hyper-parameter  $\lambda \in \{1, 10^{-1}, \dots, 10^{-6}, 0\}$ . For the paired IgT5 and IgBert models, the variable region embedding is computed as the average over both the heavy and light chain tokens, while for all other models, we compute separately an embedding representation for the heavy and light chain by computing the average over their token representations, and then concatenate the heavy and light sequence vectors into a higher-dimensional feature vector. Thus, the paired models require cross-chain features that persist after averaging, and are condensed to a feature vector of half the size of the unpaired models.

Despite this, as can be seen in Table 3, the highest  $R^2$  on the binding energy prediction tasks comes from the paired models, demonstrating the importance of training on natively paired

**Table 3.**  $R^2$  for a linear model applied on the embeddings of each language model to predict binding energy to an antigen or expression. The best, second and third best performing models for each benchmark are shown in bold, underlined and italic respectively.

Model	Binding	Binding	Binding	Expression
	$N = 422$	$N = 2048$	$N = 4275$	$N = 4275$
	[42]	[43]	[44]	[44]
AbLang [5]	$0.293 \pm 0.117$	<u><math>0.246 \pm 0.038</math></u>	<u><math>0.244 \pm 0.034</math></u>	$0.439 \pm 0.027$
AntiBERTy [26]	$0.239 \pm 0.102$	$0.217 \pm 0.056$	$0.199 \pm 0.025$	$0.401 \pm 0.032$
ProtBert [20]	$0.200 \pm 0.106$	$0.149 \pm 0.024$	$0.101 \pm 0.017$	$0.491 \pm 0.029$
IgBert-unpaired	$0.278 \pm 0.094$	$0.181 \pm 0.040$	$0.177 \pm 0.018$	$0.347 \pm 0.023$
IgBert	<b><math>0.306 \pm 0.114</math></b>	$0.131 \pm 0.047$	$0.174 \pm 0.032$	$0.400 \pm 0.023$
ProtT5 [20]	$0.290 \pm 0.105$	$0.186 \pm 0.037$	$0.206 \pm 0.029$	<b><math>0.697 \pm 0.02</math></b>
IgT5-unpaired	<u><math>0.299 \pm 0.119</math></u>	<u><math>0.245 \pm 0.049</math></u>	$0.179 \pm 0.014$	$0.567 \pm 0.025$
IgT5	$0.274 \pm 0.070$	<b><math>0.297 \pm 0.057</math></b>	<b><math>0.25 \pm 0.019</math></b>	$0.548 \pm 0.067$

<https://doi.org/10.1371/journal.pcbi.1012646.t003>

data, which was also recently highlighted by [27]. The general protein models, ProtBert and ProtT5, perform below the antibody-specific ones on binding prediction tasks. However, on the expression task, the general protein language models outperform the antibody specific models, with ProtT5 having the highest  $R^2$ . This suggests that evolutionary information or broader patterns across different protein families present in general protein models are important in general property downstream tasks and are not learned as effectively by antibody language models.

We use the  $R^2$  metric due to its robustness in evaluating the predictive accuracy of our model across diverse datasets. A similar analysis showing the Pearson correlation can be found in [S4 Appendix](#).

## Discussion

In this study, we investigated whether large antibody-specific language models could improve performance on sequence recovery and downstream predictive tasks.

We pre-trained both T5 and BERT models on over two billion unpaired sequences from the Observed Antibody Space using an MLM objective. We then fine-tuned the unpaired models on paired heavy and light sequences to create paired models that can learn cross-chain features. The resulting IgBert and IgT5 models are made available through Zenodo [17] and Huggingface. Models available at [huggingface.co/Exscientia](https://huggingface.co/Exscientia).

We identify several key takeaways that shed light on the potential and limitations of these specialised language models.

One striking result was the substantial improvement in performance achieved through fine-tuning our models on paired data. This improvement can be attributed to the models' ability to learn cross-chain features, facilitating a deeper understanding of antibody sequences. Furthermore, we believe that the quality of the paired data played a pivotal role in enhancing performance, highlighting the significance of data quality in training protein language models. Experimental verification of designs derived from IgT5 and IgBert would provide an interesting avenue for validating this performance. It would also be valuable to study the scaling of model performance with the amount of paired sequence data used in training, and whether additional sequencing data is likely to lead to substantial improvements in down stream tasks.

While our experiments demonstrated improved binding affinity predictions using embeddings from large antibody-specific language models, we also observed a discrepancy in their ability to predict expression compared to general protein models. This result suggests that general protein models, which have acquired knowledge about the evolutionary diversity of proteins, have an advantage in capturing patterns relating to protein expression, whilst antibody language models may be better suited at learning specialised features relevant to antibody specific properties.

The results presented in this article open the avenue for several promising future directions. A second stage pre-training approach that incorporates general protein sequences into the training data would allow the model to enhance generalisability, improving its performance on expression benchmarks. Additionally, the integration of structural and functional information into the language model embedding representations through the application of methods such as contrastive learning [45] might improve performance on property prediction tasks, and there has been promising recent work in this direction [46, 47]. It would also be interesting to consider more parameter efficient fine-tuning strategies such as low rank adaptation [48], and training strategies that can maintain expertise in previous tasks such as elastic weight consolidation [49] to enable more scalable and robust specialised protein language models. More extended ablation studies of the impact on model performance of

clustering thresholds used in preparing the training and test data would be valuable, as these have been shown to impact training dynamics in general protein language models [50]. Notably, it would be useful to understand whether different clustering choices, e.g. clustering along the CDRs instead of the variable region sequence, or training on cluster representatives instead of all sequences, impact training dynamics at scale for different sequence identity thresholds. The models trained in this study might provide a particularly powerful approach when exploring more intricate downstream tasks, such as structure prediction. Furthermore, the paired models offer an interesting opportunity to leverage learned cross-chain features to predict pairing relationships. These models could also be leveraged for *in silico* affinity maturation, incorporating them into a therapeutic candidate generation pipelines to improve developability [51, 52].

Extending the capabilities of our IgT5 model by fine-tuning it for generative tasks also opens up opportunities for generating novel antibody sequences that can provide a powerful approach for antibody design.

Our research on antibody-specific language models highlights their potential for advancing antibody-related research and development. The work presented in this article provides a key step towards harnessing the power of these models for therapeutic development, and paves the way for breakthroughs on the application of specialised language models for antibody engineering and drug discovery.

## Supporting information

**S1 Appendix. Model usage.** Example code for the use of the IgT5(-unpaired) and IgBert (-unpaired) models, along with links to more detailed information.  
(PDF)

**S2 Appendix. OAS dataset.** Overview of the training dataset, which consists of the OAS database accessed in November 2023.  
(PDF)

**S3 Appendix. Unpaired sequence recovery.** The sequence recovery achieved by each model on a test set of unpaired sequences.  
(PDF)

**S4 Appendix. Correlation for binding affinity and expression.** The correlation of the predicted binding energy and expression with the data.  
(PDF)

## Acknowledgments

We are grateful to the NVIDIA Corporation for providing access to the 4 DGX nodes used in this study. We thank Daniel Cutting and Constantin Schneider for useful discussions.

## Author Contributions

**Conceptualization:** Henry Kenlay, Frédéric A. Dreyer, Charlotte M. Deane.

**Data curation:** Frédéric A. Dreyer, Aleksandr Kovaltsuk.

**Formal analysis:** Henry Kenlay.

**Investigation:** Henry Kenlay, Frédéric A. Dreyer.

**Methodology:** Henry Kenlay, Frédéric A. Dreyer.

**Software:** Henry Kenlay, Frédéric A. Dreyer, Dom Miketa.

**Supervision:** Douglas Pires, Charlotte M. Deane.

**Validation:** Henry Kenlay, Frédéric A. Dreyer.

**Visualization:** Frédéric A. Dreyer.

**Writing – original draft:** Frédéric A. Dreyer.

**Writing – review & editing:** Henry Kenlay, Frédéric A. Dreyer, Aleksandr Kovaltsuk, Dom Miketa, Douglas Pires, Charlotte M. Deane.

## References

1. Narciso JE, Uy I, Cabang A, Chavez J, Pablo J, Padilla-Concepcion G, et al. Analysis of the antibody structure based on high-resolution crystallographic studies. *New biotechnology*. 2011; 28:435–47. <https://doi.org/10.1016/j.nbt.2011.03.012> PMID: 21477671
2. Tsuchiya Y, Mizuguchi K. The diversity of H3 loops determines the antigen-binding tendencies of anti-body CDR loops. *Protein science: a publication of the Protein Society*. 2016; 25. <https://doi.org/10.1002/pro.2874> PMID: 26749247
3. Regep C, Georges G, Shi J, Popovic B, Deane C. The H3 loop of antibodies shows unique structural characteristics: CDR H3 of Antibodies Shows Unique Structural Characteristics. *Proteins: Structure, Function, and Bioinformatics*. 2017; 85. <https://doi.org/10.1002/prot.25291> PMID: 28342222
4. Chaudhary N, Wesemann DR. Analyzing Immunoglobulin Repertoires. *Frontiers in Immunology*. 2018; 9. <https://doi.org/10.3389/fimmu.2018.00462> PMID: 29593723
5. Olsen TH, Moal IH, Deane CM. AbLang: an antibody language model for completing antibody sequences. *Bioinformatics Advances*. 2022; 2(1):vbac046. <https://doi.org/10.1093/bioadv/vbac046> PMID: 36699403
6. Shuai RW, Ruffolo JA, Gray JJ. Generative language modeling for antibody design. *bioRxiv*. 2022.
7. Gruver N, Stanton S, Frey NC, Rudner TGJ, Hotzel I, Lafrance-Vanassee J, et al. Protein Design with Guided Discrete Diffusion; 2023.
8. Ruffolo JA, Chu LS, Mahajan SP, Gray JJ. Fast, accurate antibody structure prediction from deep learning on massive set of natural antibodies. *Nature Communications*. 2023; 14(1):2389. <https://doi.org/10.1038/s41467-023-38063-x> PMID: 37185622
9. Høie M, Hummer A, Olsen T, Nielsen M, Deane C. AntiFold: Improved antibody structure design using inverse folding. In: NeurIPS 2023 Generative AI and Biology (GenBio) Workshop; 2023. Available from: <https://openreview.net/forum?id=bxZMKHtlL6>.
10. Leem J, Mitchell LS, Farmery JHR, Barton J, Galson JD. Deciphering the language of antibodies using self-supervised learning. *Patterns*. 2022; 3(7):100513. <https://doi.org/10.1016/j.patter.2022.100513> PMID: 35845836
11. Prihoda D, Maamary J, Waight A, Juan V, Fayadat-Dilman L, Svozi D, et al. BioPhi: A platform for antibody design, humanization, and humanness evaluation based on natural antibody repertoires and deep learning. *mAbs*. 2022; 14(1):2020203. <https://doi.org/10.1080/19420862.2021.2020203> PMID: 35133949
12. Harmalkar A, Rao R, Xie YR, Honer J, Deisting W, Anlahr J, et al. Toward generalizable prediction of antibody thermostability using machine learning on sequence and structure features. *mAbs*. 2023; 15(1):2163584. <https://doi.org/10.1080/19420862.2022.2163584> PMID: 36683173
13. Kovaltsuk A, Leem J, Kelm S, Snowden J, Deane CM, Krawczyk K. Observed Antibody Space: A Resource for Data Mining Next-Generation Sequencing of Antibody Repertoires. *The Journal of Immunology*. 2018; 201(8):2502–2509. <https://doi.org/10.4049/jimmunol.1800708> PMID: 30217829
14. Olsen TH, Boyles F, Deane CM. Observed Antibody Space: A diverse database of cleaned, annotated, and translated unpaired and paired antibody sequences. *Protein Science*. 2022; 31(1):141–146. <https://doi.org/10.1002/pro.4205> PMID: 34655133
15. Devlin J, Chang M, Lee K, Toutanova K. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *CoRR*. 2018;abs/1810.04805.
16. Raffel C, Shazeer N, Roberts A, Lee K, Narang S, Matena M, et al. Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer. *J Mach Learn Res*. 2020; 21(1).
17. Dreyer FA, Kenlay H. Large scale paired antibody language models; 2024. Available from: <https://zenodo.org/doi/10.5281/zenodo.10876908>.

18. Rives A, Meier J, Sercu T, Goyal S, Lin Z, Liu J, et al. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences*. 2021; 118(15):e2016239118. <https://doi.org/10.1073/pnas.2016239118> PMID: 33876751
19. Lin Z, Akin H, Rao R, Hie B, Zhu Z, Lu W, et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*. 2023; 379(6637):1123–1130. <https://doi.org/10.1126/science.adc2574> PMID: 36927031
20. Elnaggar A, Heinzinger M, Dallago C, Rehwald G, Wang Y, Jones L, et al. ProtTrans: Toward Understanding the Language of Life Through Self-Supervised Learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 2022; 44(10):7112–7127. <https://doi.org/10.1109/TPAMI.2021.3095381> PMID: 34232869
21. Madani A, Krause B, Greene ER, Subramanian S, Mohr BP, Holton JM, et al. Large language models generate functional protein sequences across diverse families. *Nature Biotechnology*. 2023; 41(8):1099–1106. <https://doi.org/10.1038/s41587-022-01618-2> PMID: 36702895
22. Consortium TU. UniProt: the Universal Protein Knowledgebase in 2023. *Nucleic Acids Research*. 2022; 51(D1):D523–D531. <https://doi.org/10.1093/nar/gkac1052>
23. Bairoch A, Apweiler R. The SWISS-PROT protein sequence database and its supplement TrEMBL in 2000. *Nucleic Acids Research*. 2000; 28(1):45–48. <https://doi.org/10.1093/nar/28.1.45> PMID: 10592178
24. Hie BL, Yang KK, Kim PS. Evolutionary velocity with protein language models predicts evolutionary dynamics of diverse proteins. *Cell Systems*. 2022; 13(4):274–285. <https://doi.org/10.1016/j.cels.2022.01.003> PMID: 35120643
25. Liu Y, Ott M, Goyal N, Du J, Joshi M, Chen D, et al. RoBERTa: A Robustly Optimized BERT Pretraining Approach; 2019.
26. Ruffolo JA, Gray JJ, Sulam J. Deciphering antibody affinity maturation with language models and weakly supervised learning; 2021.
27. Burbach SM, Briney B. Improving antibody language models with native pairing; 2023.
28. Olsen TH, Moal IH, Deane CM. Addressing the antibody germline bias and its effect on language models for improved antibody design. *bioRxiv*. 2024. <https://doi.org/10.1101/bioinformatics/btae618> PMID: 39460949
29. Vashchenko D, Nguyen S, Goncalves A, da Silva FL, Petersen B, Desautels T, et al. AbBERT: Learning Antibody Humanness via Masked Language Modeling. *bioRxiv*. 2022.
30. Chu SKS, Wei KY. Generative Antibody Design for Complementary Chain Pairing Sequences through Encoder-Decoder Language Model; 2023.
31. Chen B, Cheng X, ao Geng Y, Li S, Zeng X, Wang B, et al. xTrimoPGLM: Unified 100B-Scale Pre-trained Transformer for Deciphering the Language of Protein. *bioRxiv*. 2023.
32. Vaswani A, Shazeer N, Parmar N, Uszkoreit J, Jones L, Gomez AN, et al. Attention Is All You Need; 2017.
33. Suzek BE, Wang Y, Huang H, McGarvey PB, Wu CH, the UniProt Consortium. UniRef clusters: a comprehensive and scalable alternative for improving sequence similarity searches. *Bioinformatics*. 2014; 31(6):926–932. <https://doi.org/10.1093/bioinformatics/btu739> PMID: 25398609
34. Steinegger M, Mirdita M, Söding J. Protein-level assembly increases protein sequence recovery from metagenomic samples manyfold. *Nature Methods*. 2019; 16(7):603–606. <https://doi.org/10.1038/s41592-019-0437-4> PMID: 31235882
35. Steinegger M, Söding J. Clustering huge protein sequence sets in linear time. *Nature Communications*. 2018; 9(1):2542. <https://doi.org/10.1038/s41467-018-04964-5> PMID: 29959318
36. Steinegger M, Söding J. MMseqs2 enables sensitive protein sequence searching for the analysis of massive data sets. *Nature Biotechnology*. 2017; 35(11):1026–1028. <https://doi.org/10.1038/nbt.3988> PMID: 29035372
37. Rajbhandari S, Rasley J, Ruwase O, He Y. Zero: Memory optimizations toward training trillion parameter models. In: SC20: International Conference for High Performance Computing, Networking, Storage and Analysis. IEEE; 2020. p. 1–16.
38. Loshchilov I, Hutter F. Decoupled Weight Decay Regularization. In: International Conference on Learning Representations; 2019. Available from: <https://openreview.net/forum?id=Bkg6RiCqY7>.
39. Kirkpatrick J, Pascanu R, Rabinowitz N, Veness J, Desjardins G, Rusu AA, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*. 2017; 114(13):3521–3526. <https://doi.org/10.1073/pnas.1611835114> PMID: 28292907
40. Chen S, Hou Y, Cui Y, Che W, Liu T, Yu X. Recall and Learn: Fine-tuning Deep Pretrained Language Models with Less Forgetting; 2020.

41. Chungyoun M, Ruffolo J, Gray JJ. FLaB: Benchmarking tasks in fitness landscape inference for antibodies. NeurIPS Machine Learning in Structural Biology Workshop. 2023;.
42. Shanehsazzadeh A, Bachas S, Kasun G, Sutton JM, Steiger AK, Shuai R, et al. Unlocking de novo antibody design with generative artificial intelligence. bioRxiv. 2023.
43. Warszawski S, Borenstein Katz A, Lipsh R, Khmelnitsky L, Ben Nissan G, Javitt G, et al. Optimizing antibody affinity and stability by the automated design of the variable light-heavy chain interfaces. PLOS Computational Biology. 2019; 15(8):1–24. <https://doi.org/10.1371/journal.pcbi.1007207> PMID: 31442220
44. Koenig P, Lee CV, Walters BT, Janakiraman V, Stinson J, Patapoff TW, et al. Mutational landscape of antibody variable domains reveals a switch modulating the interdomain conformational dynamics and antigen binding. Proceedings of the National Academy of Sciences. 2017; 114(4):E486–E495. <https://doi.org/10.1073/pnas.1613231114> PMID: 28057863
45. Gao T, Yao X, Chen D. SimCSE: Simple Contrastive Learning of Sentence Embeddings; 2022.
46. Hamamsy T, Barot M, Morton JT, Steinegger M, Bonneau R, Cho K. Learning sequence, structure, and function representations of proteins with language models. bioRxiv. 2023. <https://doi.org/10.1101/2023.11.26.568742> PMID: 38045331
47. Barton J, Galson J, Leem J. Enhancing Antibody Language Models with Structural Information. NeurIPS Machine Learning in Structural Biology workshop. 2023;.
48. Hu EJ, Shen Y, Wallis P, Allen-Zhu Z, Li Y, Wang S, et al. LoRA: Low-Rank Adaptation of Large Language Models. CoRR. 2021;abs/2106.09685.
49. Kirkpatrick J, Pascanu R, Rabinowitz NC, Veness J, Desjardins G, Rusu AA, et al. Overcoming catastrophic forgetting in neural networks. CoRR. 2016;abs/1612.00796.
50. Meier J, Rao R, Verkuil R, Liu J, Sercu T, Rives A. Language models enable zero-shot prediction of the effects of mutations on protein function. In: Ranzato M, Beygelzimer A, Dauphin Y, Liang PS, Vaughan JW, editors. Advances in Neural Information Processing Systems. vol. 34. Curran Associates, Inc.; 2021. p. 29287–29303. Available from: [https://proceedings.neurips.cc/paper\\_files/paper/2021/file/f51338d736f95dd42427296047067694-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2021/file/f51338d736f95dd42427296047067694-Paper.pdf).
51. Hie BL, Shanker VR, Xu D, Bruun TUJ, Weidenbacher PA, Tang S, et al. Efficient evolution of human antibodies from general protein language models. Nature Biotechnology. 2023. <https://doi.org/10.1038/s41587-023-01763-2> PMID: 37095349
52. Outeiral C, Deane CM. Perfecting antibodies with language models. Nature Biotechnology. 2023.